
Deprivelege Root access of a daemon

By:

Ram Jangir (CS21M517)
Sneha M R (CS21M522)

What is sysvinit?

- It is the first process (init process pid=0) started by the kernel when you boot up any Linux or Unix system. It means other processes are its child in one or the other way.
- Sysvinit process continues to run and waits for special commands like ‘shutdown’, which are used to shut down a Linux system.

Disadvantages:

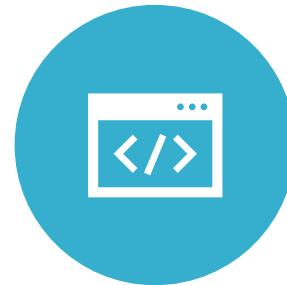
- Sysvinit remained a perfect system to bring up and shutdown Linux-based systems. But as time passed by, the system became slow and inflexible, especially for modern-day computers/systems.

In 2010 systemd was proposed to replace the widely used sysvinit system.

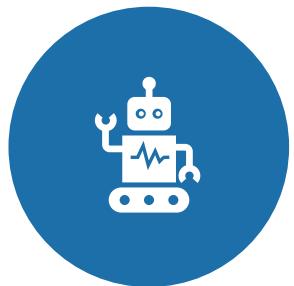
Advantages of systemd over sysvinit:



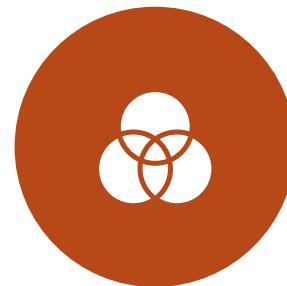
Services are started in parallel instead of serially, which reduces boot-up time.



Service configuration is easier through service files.



It has more robust control and debugging like journalctl.



It is compatible with **Sysvinit**.

Service File Options

How to Improve:

- **Service file Options:**
 - User= and Group= options in systemd service files
 - Allows launching of daemons as non-root users/groups without needing setuid/setgid capabilities.
- **Systemd-tmpfiles**
 - Daemon that runs on startup
 - Reads configuration files supplied by /etc/tmpfiles.d/*/*.conf and creates/modifies files after bootup.
 - Runs as root so your daemon doesn't have to.
- **Udev**
 - Continuously running daemon
 - Updates sysfs permissions once at bootup and monitors uevents for any dynamically created nodes from the kernel. If new nodes are added – this daemon will update the permissions

Example service file

```
cat /etc/systemd/system/app.service
```

```
[Unit]
```

```
SourcePath=/local/mnt/workspace/SERVICE_FILE/app  
Description= APP Service file
```

```
[Service]
```

```
User=radio  
Group=radio
```

```
ExecStart=/local/mnt/workspace/SERVICE_FILE/app  
ExecStop= /usr/bin/kill -9 app
```

```
SupplementaryGroups=diag inet net_admin system wifi netdev net_raw kmsq
```

```
CapabilityBoundingSet=CAP_NET_ADMIN CAP_NET_RAW CAP_NET_BIND_SERVICE CAP_SETUID CAP_SETGID
```

```
AmbientCapabilities=CAP_NET_ADMIN CAP_NET_RAW CAP_NET_BIND_SERVICE CAP_SETUID CAP_SETGID
```

Detailed service file:

[Unit]

SourcePath=/local/mnt/workspace/SERVICE_FILE/app

Description= APP Service file

[Service]

User=radio

Group=radio

ExecStart=/local/mnt/workspace/SERVICE_FILE/app

ExecStop= /usr/bin/kill -9 app

RemainAfterExit=

Type=

Wants=

After=

Before=

Part of=

WantedBy=

RequiredBy=

Restart-on-failure=yes

SupplementaryGroups=diag inet net_admin system wifi netdev net_raw kmsq

CapabilityBoundingSet=CAP_NET_ADMIN CAP_NET_RAW CAP_NET_BIND_SERVICE CAP_SETUID CAP_SETGID

AmbientCapabilities=CAP_NET_ADMIN CAP_NET_RAW CAP_NET_BIND_SERVICE CAP_SETUID CAP_SETGID

[Install]

Wantedby=multi-user.target

RUNLEVEL (RHEL 6)	Target (RHEL 7)	Description
0	poweroff.target	To Halt/Shutdown the System
1	rescue.target	To Start in Single User Mode (For Troubleshooting/Administration Tasks)
2	multi-user.target	System Starts with Multiuser Mode but without Networking like (NFS)
3	multi-user.target	With Multiuser Mode with Networking
4	multi-user.target	Reserved
5	graphical.target	With Graphical User Mode (GUI)
6	reboot.target	Reboot the System

Systemd Commands:

- `systemctl start app` # start the service. (app is the service file name)
- `systemctl stop app` # Stop the service.
- `systemctl restart app` # Restart the service.
- `systemctl status app` # Get the status of the service
- `journalctl` # Get all the logs
- `journalctl -u app` # Get the logs of the particular service
- `journalctl -u app -n 10` # Get the logs of the particular service, **10 lines**
- `systemctl enable app` #Run everytime after bootup
- `systemctl disable app` #Dont run after bootup.

Challenge 1)

- *Procfs entries and network related commands are impacted:*
- */proc/sys/net/ipv6/conf/disable_ipv6*
- */proc/sys/net/ipv4/ip_forward*
- */proc/sys/net/netfilter/nf_conntrack_udp_timeout_stream*
- */proc/sys/debug/sfe/packet_stats_on*
- *Brctl addif wlan0 bridge0*
- *iw dev wlan0 set 4addr on*
- *Ifconfig wlan0 up, down*
- *Vconfig <iface name>*
- *Solution:*
- *Giving the NET_ADMIN and NET_RAW capabilities*

Systemd-tmpfiles

How to Improve:

- Service file Options:
 - User= and Group= options in systemd service files
 - Allows launching of daemons as non-root users/groups without needing setuid/setgid capabilities.
- **Systemd-tmpfiles:**
 - Daemon that runs on startup
 - Reads configuration files supplied by `/etc/tmpfiles.d/*.conf` and creates/modifies files after bootup.
 - Runs as root so your daemon doesn't have to.
- Udev
 - Continuously running daemon
 - Updates sysfs permissions once at bootup and monitors uevents for any dynamically created nodes from the kernel. If new nodes are added – this daemon will update the permissions

Systemd-tmpfiles:

```
$cat /etc/tmpfiles.d/test.conf
```

```
f /local/mnt/workspace/tmp.txt 0755 smaganah nobody --  
d /local/mnt/workspace/tmp_dir 0755 smaganah nobody --  
L /local/mnt/workspace/tmp_link 0755 smaganah nobody --
```

Command:

```
$systemd-tmpfiles --create
```

Output:

```
root@blr-ubuntu-smaganah:/local/mnt/workspace# ls -lrt | grep -i tmp  
-rwxr-xr-x 1 smaganah nobody 0 Nov 29 10:28 tmp.txt  
drwxr-xr-x 2 smaganah nobody 4096 Nov 29 10:28 tmp_dir  
lrwxrwxrwx 1 root root 48 Nov 29 10:28 tmp_link -> /usr/share/factory//local/mnt/workspace/tmp_link  
root@blr-ubuntu-smaganah:/local/mnt/workspace#
```

Challenge 2)

- /var/run and /tmp are impacted:
 - /tmp/statefdbtable.txt
 - /tmp/session_token
 - /tmp/brctl.tmp
 - /var/run/dnsmasq_socket
- *Solution:*
- *Create /var/run/data and /tmp/data , the data folders have radio:radio permission. Hence only Private daemons can access it. Use systemd-tmpfiles to create these folders.*

Challenge 3)

/etc/data is impacted:

- *Solution:*
- *The config files used by the private daemons is present in /etc/data should be changed to radio:radio permission.*
- *The config file which are used by both private daemon and opensource daemon (eg: dnsmasq.conf) should be changed to root:radio permission so that both Private-daemon (runs as radio) and open-source daemon (runs as root) should be able to edit it.*

Udev events

How to Improve:

- Service file Options:
 - User= and Group= options in systemd service files
 - Allows launching of daemons as non-root users/groups without needing setuid/setgid capabilities.
- **systemd-tmpfiles**
 - Daemon that runs on startup
 - Reads configuration files supplied by /etc/tmpfiles.d/* .conf and creates/modifies files after bootup.
 - Runs as root so your daemon doesn't have to.
- **Udev:**
 - Continuously running daemon
 - Updates sysfs permissions once at bootup and monitors uevents for any dynamically created nodes from the kernel. If new nodes are added – this daemon will update the permissions

- **udevadm info /dev/tty0**

```
smaganah@blr-ubuntu-smaganah:/local/mnt/workspace$ udevadm info /dev/tty0
P: /devices/virtual/tty/tty0
N: tty0
E: DEVNAME=/dev/tty0
E: DEVPATH=/devices/virtual/tty/tty0
E: ID_MM_CANDIDATE=1
E: MAJOR=4
E: MINOR=0
E: SUBSYSTEM=tty
E: USEC_INITIALIZED=15858153
smaganah@blr-ubuntu-smaganah:/local/mnt/workspace$ smaganah@blr-ubuntu-smaganah:/local/mnt/workspace$
```

Commands to trigger Udev event:
udevadm info /dev/tty0
udevadm control --reload-rules
udevadm trigger

\$cat /etc/udev/rules.d/udev_rules.rules

```
SUBSYSTEM=="tty" RUN+=" /usr/bin/udev_script.sh"
```

\$cat /usr/bin/udev_script.sh

```
#!/bin/sh
```

```
chown -h smaganah.nobody /dev/tty0
```

Output:

```
crw--w----. 1 smaganah nobody 4, 0 Nov 30 16:02 /dev/tty0
```

Challenge 4)

Printing logs to /dev/kmsg is impacted:

Solution:

/var/log/messages should be used instead of /dev/kmsg

/dev/ttyX is impacted:

Solution:

/dev/ttyX permission should be changed using udev rules.

System DBUS (SDBus)

System Dbus:

- Sdbus allows depivileged process to start a process in root:root (user:group)

Example:

DBUS API to start the daemon:

```
r = sd_bus_call_method(bus,
                        "org.freedesktop.systemd1",
                        "/org/freedesktop/systemd1",
                        "org.freedesktop.systemd1.Manager",
                        "StartUnit", //specify to start the process. Similary StopUnit, RestartUnit
                        &error,
                        &m,
                        "ss",
                        "tinyproxy.service", //specify which service to start or stop
                        "replace");
```

It will start the tinyproxy service.

Sample Open source daemon service file:

\$cat /etc/init.d/tinyproxy.service

```
[Unit]
Description=dnsmasq Service
SourcePath=/usr/sbin/tinyproxy

[Service]
User=nobody
Group=inet
Restart=no
Type=oneshot
RemainAfterExit=yes
ExecStart=/usr/sbin/tinyproxy
ExecStop=/usr/bin/killall -15 tinyproxy

[Install]
WantedBy=multi-user.target
```

Installation:

```
sudo apt-get install tinyproxy
sudo apt-get install libsystemd-dev
```

Sample Program:

```
1 #include <stdio.h>
2 #include <systemd/sd-bus.h>
3 #include<errno.h>
4
5 int main()
6 {
7
8     sd_bus *bus = NULL;
9     int dbus_ret;
10
11    sd_bus_message *m = NULL;
12    int ret;
13    sd_bus_error error = SD_BUS_ERROR_NULL;
14
15    /* Dbus initialization*/
16    dbus_ret = sd_bus_open_system(&bus);
17    if (dbus_ret < 0)
18    {
19        printf("sd_bus: Failed to connect to system bus: %d \n", dbus_ret);
20
21    }
22    else
23    {
24        printf("sd_bus: Successfully connected to systembus \n");
25    }
26
27
28    ret = sd_bus_call_method(bus,
29                            "org.freedesktop.systemd1",
30                            "/org/freedesktop/systemd1",
31                            "org.freedesktop.systemd1.Manager",
32                            "RestartUnit",
33                            &error,
34                            &m,
35                            "ss",
36                            "tinyproxy.service",
37                            "replace");
38
39    if (ret < 0)
40    {
41        printf("SERVICE START FAILED error.msg:%s ret:%d errno:%d\n",error.message, ret, errno);
42    } else {
43
44        printf ("SERVICE STARTED SUCCESSFULLY %d %d", ret, errno);
45    }
46
47
48
49    return 0;
50 }
```

Compile:

```
$gcc sd_bus.c -lsystemd
```

Output:

```
$./a.out
```

```
sd_bus: Successfully connected to systembus
SERVICE STARTED SUCCESSFULLY 1 11
```

\$ps -ef | grep -i tinyproxy

```
nobody    7923      1  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7936    7923  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7937    7936  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7938    7937  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7939    7938  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7940    7939  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7941    7940  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7942    7941  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7943    7942  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7944    7943  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
nobody    7945    7944  0 21:11 ?          00:00:00 /usr/sbin/tinyproxy
```

\$systemctl status tinyproxy

```
● tinyproxy.service - LSB: Tinyproxy HTTP proxy
   Loaded: loaded (/etc/init.d/tinyproxy; bad; vendor preset: enabled)
   Active: active (running) since Wed 2022-11-30 21:11:50 IST; 15s ago
     Docs: man:systemd-sysv-generator(8)
 Process: 7881 ExecStop=/etc/init.d/tinyproxy stop (code=exited, status=0/SUCCESS)
 Process: 7894 ExecStart=/etc/init.d/tinyproxy start (code=exited, status=0/SUCCESS)
   Tasks: 11
  Memory: 6.2M
    CPU: 38ms
  CGroup: /system.slice/tinyproxy.service
```

```
    └─7923 /usr/sbin/tinyproxy
        ├─7936 /usr/sbin/tinyproxy
        ├─7937 /usr/sbin/tinyproxy
        ├─7938 /usr/sbin/tinyproxy
        ├─7939 /usr/sbin/tinyproxy
        ├─7940 /usr/sbin/tinyproxy
        ├─7941 /usr/sbin/tinyproxy
        ├─7942 /usr/sbin/tinyproxy
        ├─7943 /usr/sbin/tinyproxy
        ├─7944 /usr/sbin/tinyproxy
        ├─7945 /usr/sbin/tinyproxy
```

Debugging with strace:

Usage:

Strace <binary name>

Eg: strace tinyproxy

Future Work:

- Implement SELinux

Thank You